

✓ 1. Import The Data

```
from google.colab import files
import io
import pandas as pd
```

```
# Upload file
uploaded = files.upload()
```



Choose Files Insulin_Dataset.csv

- **Insulin_Dataset.csv**(text/csv) - 101336 bytes, last modified: 9/5/2025 - 100% done
Saving Insulin_Dataset.csv to Insulin_Dataset.csv

```
df = pd.read_csv(io.BytesIO(uploaded['Insulin_Dataset.csv']))
```

Check Data

```
# Check Shape And First 5 Rows
print("Shape of dataset:", df.shape)
df.head()
```



Shape of dataset: (700, 28)

| | patient_id | date | time | glucose_level | target_glucose | carbs_g | protein_g | fat_g |
|---|------------|------------|----------|---------------|----------------|---------|-----------|-------|
| 0 | 3 | 2025-01-09 | 21:08:00 | 202.5 | 100.0 | 0.0 | 18.0 | 6 |
| 1 | 1 | 2025-01-01 | 13:53:00 | 214.4 | 100.0 | 35.0 | 15.0 | 9 |
| 2 | 1 | 2025-01-01 | 20:19:00 | 238.9 | 100.0 | 0.0 | 1.0 | 12 |
| 3 | 3 | 2025-01-07 | 20:22:00 | 190.9 | 100.0 | 0.0 | 4.0 | 17 |
| 4 | 1 | 2025-01-14 | 13:59:00 | 168.8 | 100.0 | 0.0 | 16.0 | 0 |

5 rows × 28 columns

Understand The Data

```
# Dataset Info
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 700 entries, 0 to 699
Data columns (total 28 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|------------------------------|----------------|---------|
| 0 | patient_id | 700 non-null | int64 |
| 1 | date | 700 non-null | object |
| 2 | time | 700 non-null | object |
| 3 | glucose_level | 695 non-null | float64 |
| 4 | target_glucose | 696 non-null | float64 |
| 5 | carbs_g | 695 non-null | float64 |
| 6 | protein_g | 693 non-null | float64 |
| 7 | fat_g | 696 non-null | float64 |
| 8 | fiber_g | 698 non-null | float64 |
| 9 | meal_type | 699 non-null | object |
| 10 | glycemic_index | 693 non-null | float64 |
| 11 | bolus_dose_units | 695 non-null | float64 |
| 12 | bolus_type | 696 non-null | object |
| 13 | basal_rate | 695 non-null | float64 |
| 14 | insulin_on_board | 698 non-null | float64 |
| 15 | duration_of_insulin_action | 693 non-null | float64 |
| 16 | time_since_last_bolus | 698 non-null | float64 |
| 17 | dynamic_IOB | 697 non-null | float64 |
| 18 | insulin_to_carb_ratio | 694 non-null | float64 |
| 19 | correction_factor | 696 non-null | float64 |
| 20 | circadian_sensitivity_factor | 695 non-null | float64 |
| 21 | stress_level | 696 non-null | object |
| 22 | activity_level | 694 non-null | object |
| 23 | sleep_hours | 693 non-null | float64 |
| 24 | recommended_bolus_units | 696 non-null | float64 |
| 25 | age | 700 non-null | int64 |
| 26 | weight_kg | 700 non-null | float64 |
| 27 | genetic_risk | 700 non-null | int64 |

dtypes: float64(19), int64(3), object(6)

memory usage: 153.3+ KB

Missing Values Check

```
print("\nMissing Values per Column:")
```

```
print(df.isnull().sum())
```



Missing Values per Column:

| | |
|----------------------------|---|
| patient_id | 0 |
| date | 0 |
| time | 0 |
| glucose_level | 5 |
| target_glucose | 4 |
| carbs_g | 5 |
| protein_g | 7 |
| fat_g | 4 |
| fiber_g | 2 |
| meal_type | 1 |
| glycemic_index | 7 |
| bolus_dose_units | 5 |
| bolus_type | 4 |
| basal_rate | 5 |
| insulin_on_board | 2 |
| duration_of_insulin_action | 7 |
| time_since_last_bolus | 2 |
| dynamic_IOB | 3 |
| insulin_to_carb_ratio | 6 |
| correction_factor | 4 |

```

circadian_sensitivity_factor    5
stress_level                    4
activity_level                  6
sleep_hours                     7
recommended_bolus_units        4
age                             0
weight_kg                       0
genetic_risk                    0
dtype: int64

```

```

# Numerical Summary Statistics
print("\nSummary Statistics:")
display(df.describe())

```



Summary Statistics:

| | patient_id | glucose_level | target_glucose | carbs_g | protein_g | fat_g |
|--------------|------------|---------------|----------------|------------|------------|------------|
| count | 700.000000 | 695.000000 | 696.0 | 695.000000 | 693.000000 | 696.000000 |
| mean | 1.994286 | 161.299281 | 100.0 | 27.428777 | 15.431457 | 10.120690 |
| std | 0.812965 | 50.703146 | 0.0 | 33.230143 | 8.841082 | 6.077185 |
| min | 1.000000 | 70.100000 | 100.0 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 1.000000 | 118.350000 | 100.0 | 0.000000 | 8.000000 | 5.000000 |
| 50% | 2.000000 | 162.100000 | 100.0 | 0.000000 | 16.000000 | 10.000000 |
| 75% | 3.000000 | 206.450000 | 100.0 | 56.500000 | 23.000000 | 15.000000 |
| max | 3.000000 | 249.800000 | 100.0 | 100.000000 | 30.000000 | 20.000000 |

8 rows × 7 columns

```

# Unique Value In Categorical Column
print("\nUnique values in categorical columns:")
for col in ["meal_type", "bolus_type", "stress_level", "activity_level", "genetic_risk"]:
    print(f"{col}: {df[col].unique()}")

```



```

Unique values in categorical columns:
meal_type: ['dinner' 'lunch' 'snack' 'breakfast' nan]
bolus_type: ['square' 'dual-wave' 'normal' nan]
stress_level: ['high' 'low' 'medium' nan]
activity_level: ['heavy' 'none' 'light' 'moderate' nan]
genetic_risk: [1 0 2]

```

✓ 2. Cleaning The Data

```
# Missing values check again
print("Missing values before cleaning:\n", df.isnull().sum())
```

```
➞ Missing values before cleaning:
  patient_id      0
  date          0
  time          0
  glucose_level    5
  target_glucose   4
  carbs_g         5
  protein_g        7
  fat_g           4
  fiber_g          2
  meal_type        1
  glycemic_index   7
  bolus_dose_units  5
  bolus_type        4
  basal_rate        5
  insulin_on_board  2
  duration_of_insulin_action  7
  time_since_last_bolus  2
  dynamic_IOB       3
  insulin_to_carb_ratio  6
  correction_factor  4
  circadian_sensitivity_factor  5
  stress_level      4
  activity_level    6
  sleep_hours       7
  recommended_bolus_units  4
  age              0
  weight_kg         0
  genetic_risk      0
  dtype: int64
```

```
# 1. Clean Missing Value In Numerical columns (mean/median se)
num_cols = df.select_dtypes(include=['float64', 'int64']).columns
for col in num_cols:
    df[col].fillna(df[col].median(), inplace=True)
```

```
➞ /tmp/ipython-input-1618714734.py:4: FutureWarning: A value is trying to be set on a c
The behavior will change in pandas 3.0. This inplace method will never work because t

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({

    df[col].fillna(df[col].median(), inplace=True)
```

```
# 2. Clean Missing Value In Categorical columns (mode se)
cat_cols = df.select_dtypes(include=['object']).columns
for col in cat_cols:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

```
➞ /tmp/ipython-input-614862806.py:4: FutureWarning: A value is trying to be set on a cc
The behavior will change in pandas 3.0. This inplace method will never work because t

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({
```

```
df[col].fillna(df[col].mode()[0], inplace=True)
```

```
# Verify cleaning
```

```
print("\nMissing values after cleaning:\n", df.isnull().sum())
```



```
Missing values after cleaning:
```

```

patient_id      0
date            0
time            0
glucose_level   0
target_glucose  0
carbs_g         0
protein_g       0
fat_g           0
fiber_g         0
meal_type       0
glycemic_index  0
bolus_dose_units 0
bolus_type      0
basal_rate      0
insulin_on_board 0
duration_of_insulin_action 0
time_since_last_bolus 0
dynamic_IOB     0
insulin_to_carb_ratio 0
correction_factor 0
circadian_sensitivity_factor 0
stress_level    0
activity_level  0
sleep_hours     0
recommended_bolus_units 0
age             0
weight_kg       0
genetic_risk    0
dtype: int64

```

```
# Save cleaned dataset
```

```
df.to_csv("insulin_predictor_cleaned.csv", index=False)
```

```
print("\n✅ Cleaned dataset saved as 'insulin_predictor_cleaned.csv'")
```



```
✅ Cleaned dataset saved as 'insulin_predictor_cleaned.csv'
```

✓ 3. Exploratory Data Analysis (EDA)

1. Data Distribution

Histograms of all numeric features help us understand how values are spread.

Example: glucose_level distribution shows if most readings are between 100–180 mg/dL or skewed above 200.

age and weight_kg distributions give insights into the patient population profile.

2. Correlation Heatmap

The heatmap highlights linear relationships between numeric variables.

Example: carbs_g and bolus_dose_units should show a positive correlation (more carbs → higher insulin dose).

glucose_level and recommended_bolus_units will also show correlation.

3. Meal Analysis

Barplot (Total Carbs by Meal Type): Which meal (breakfast/lunch/dinner/snack) contributes the most carbs.

Scatterplot (Glucose vs Insulin by Meal Type): How insulin needs vary by meal.

Different colors for each meal make patterns easy to spot.

4. Stress & Lifestyle Impact

Boxplot (Stress vs Recommended Bolus): Shows how stress affects insulin dosage (higher stress often → higher dose).

Pie Chart (Activity Levels): Distribution of patient activity levels (none/light/moderate/heavy).

5. Key Numerical Relationships

Scatterplots:

glucose_level vs bolus_dose_units → Insulin response to blood glucose.

carbs_g vs recommended_bolus_units → Direct effect of carb intake on insulin dosage.

Histograms with KDE: Smooth distribution curves for important features (glucose, carbs, insulin doses, age, weight).

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Set Notebook Style
sns.set(style="whitegrid")
```

```
#Distribution of Numerical Features
```

```
numeric_cols = df.select_dtypes(include=['int64', 'float64']).columns
num_plots = len(numeric_cols)
```

```
plt.figure(figsize=(20, 20))
```

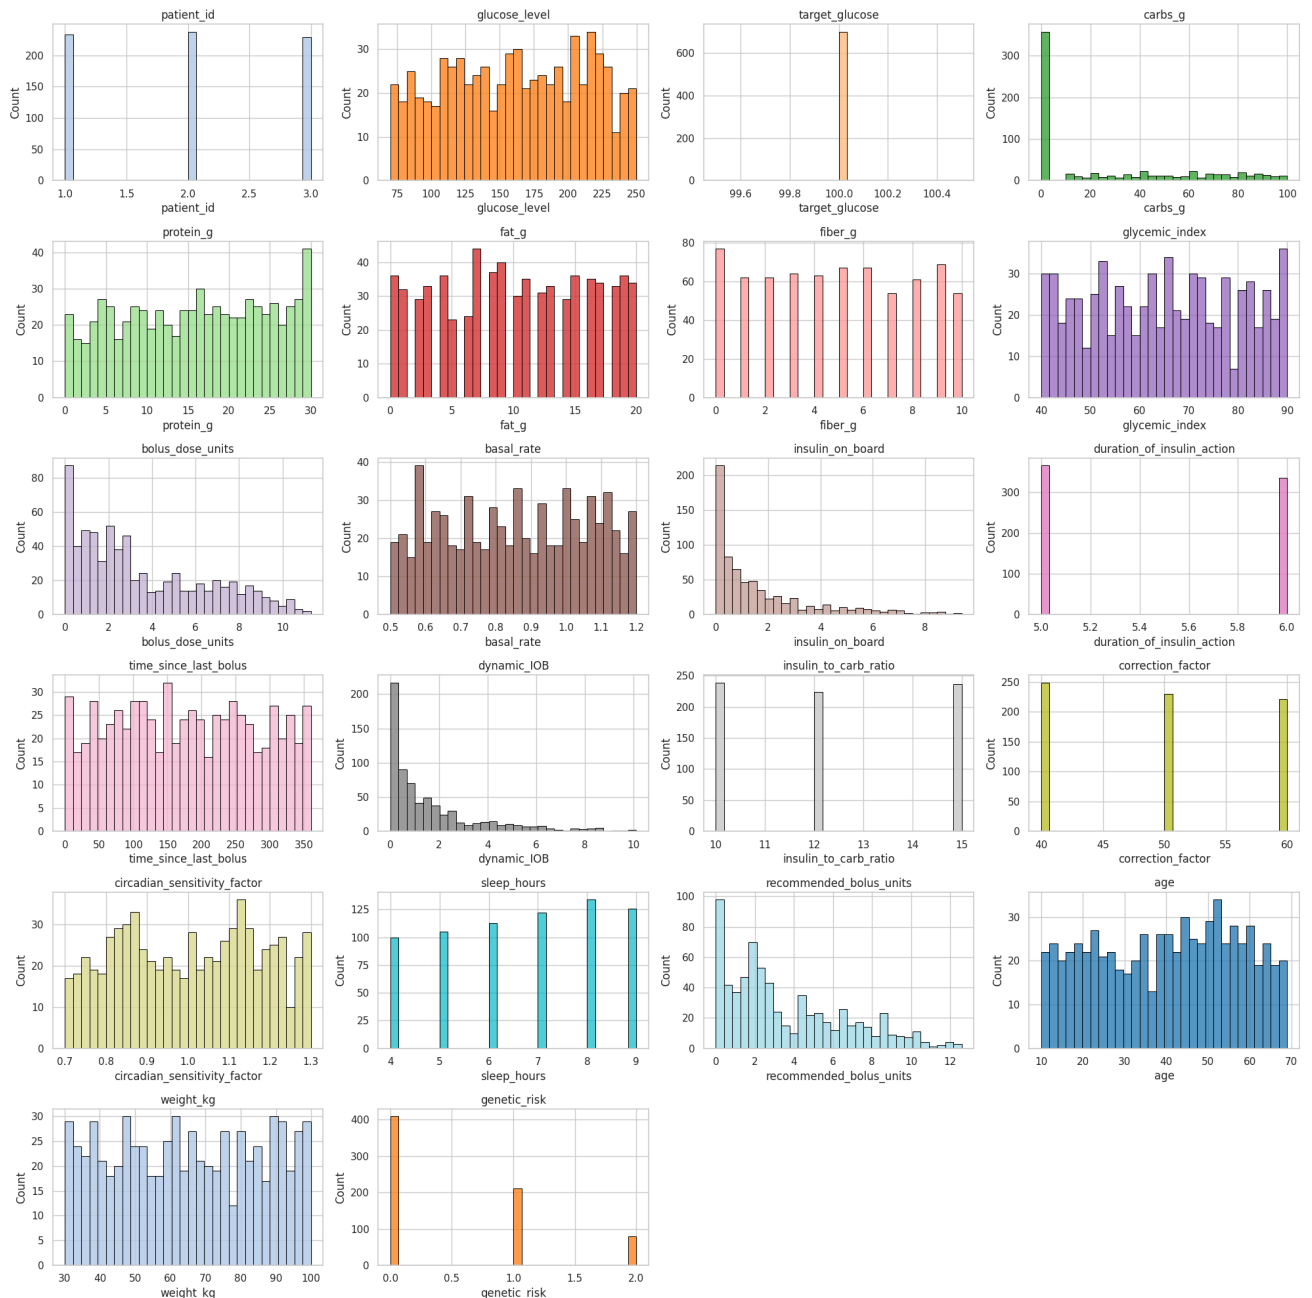
```
colors = plt.cm.tab20.colors # 20 unique colors
```

```

for i, col in enumerate(numeric_cols, 1):
    plt.subplot((num_plots//4)+1, 4, i)
    sns.histplot(df[col], bins=30, color=colors[i % len(colors)], edgecolor="black")
    plt.title(col, fontsize=12)

plt.tight_layout()
plt.show()

```



#Distribution Of Important Coloumn

```

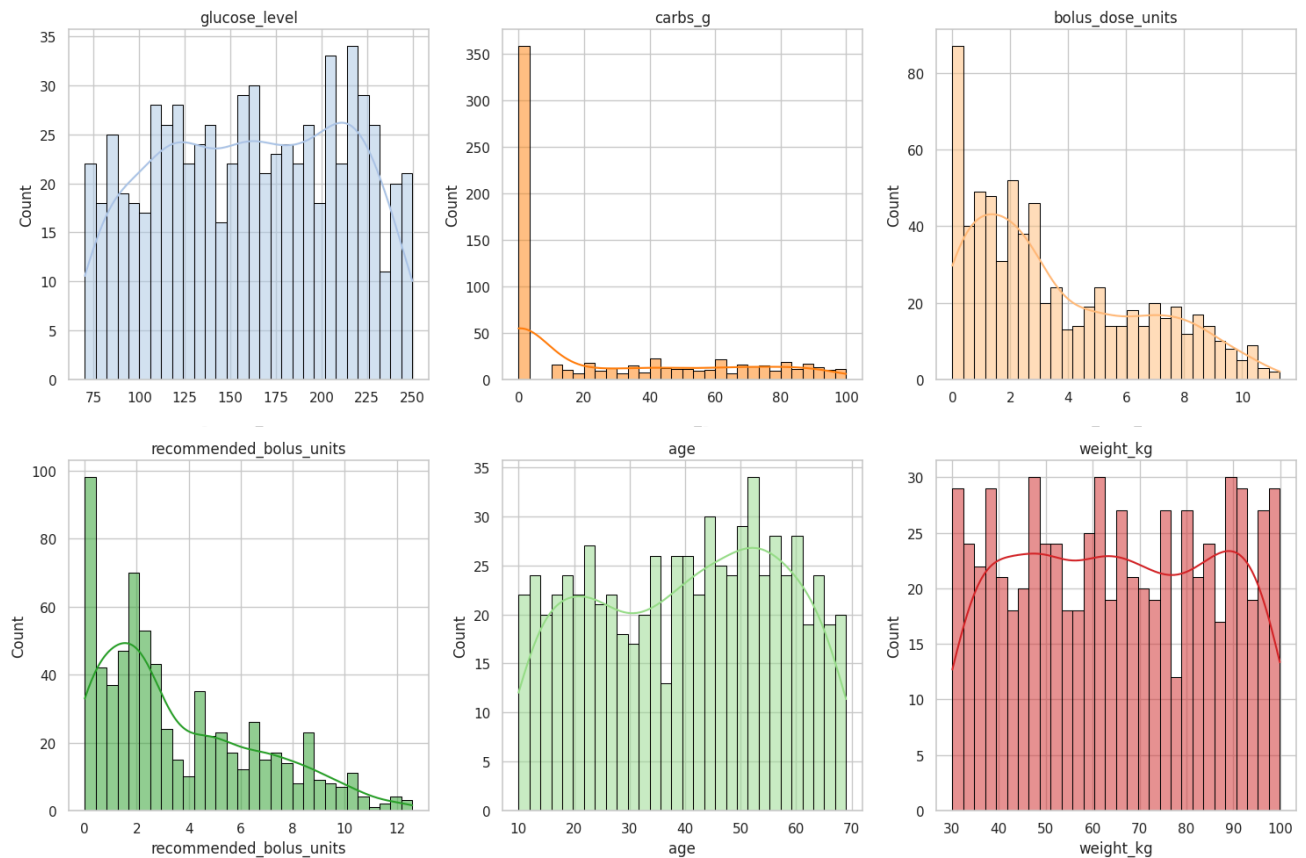
important_cols = ["glucose_level", "carbs_g", "bolus_dose_units",
                  "recommended_bolus_units", "age", "weight_kg"]

plt.figure(figsize=(15, 10))

for i, col in enumerate(important_cols, 1):
    plt.subplot(2, 3, i)
    sns.histplot(df[col], bins=30, kde=True, color=colors[i % len(colors)], edgecolor="bl
    plt.title(col, fontsize=12)

```

```
plt.tight_layout()
plt.show()
```



#Correlation

```
# Select Numeric Coloumn
```

```
numeric_df = df.select_dtypes(include=['int64', 'float64'])
```

```
# Correlation
```

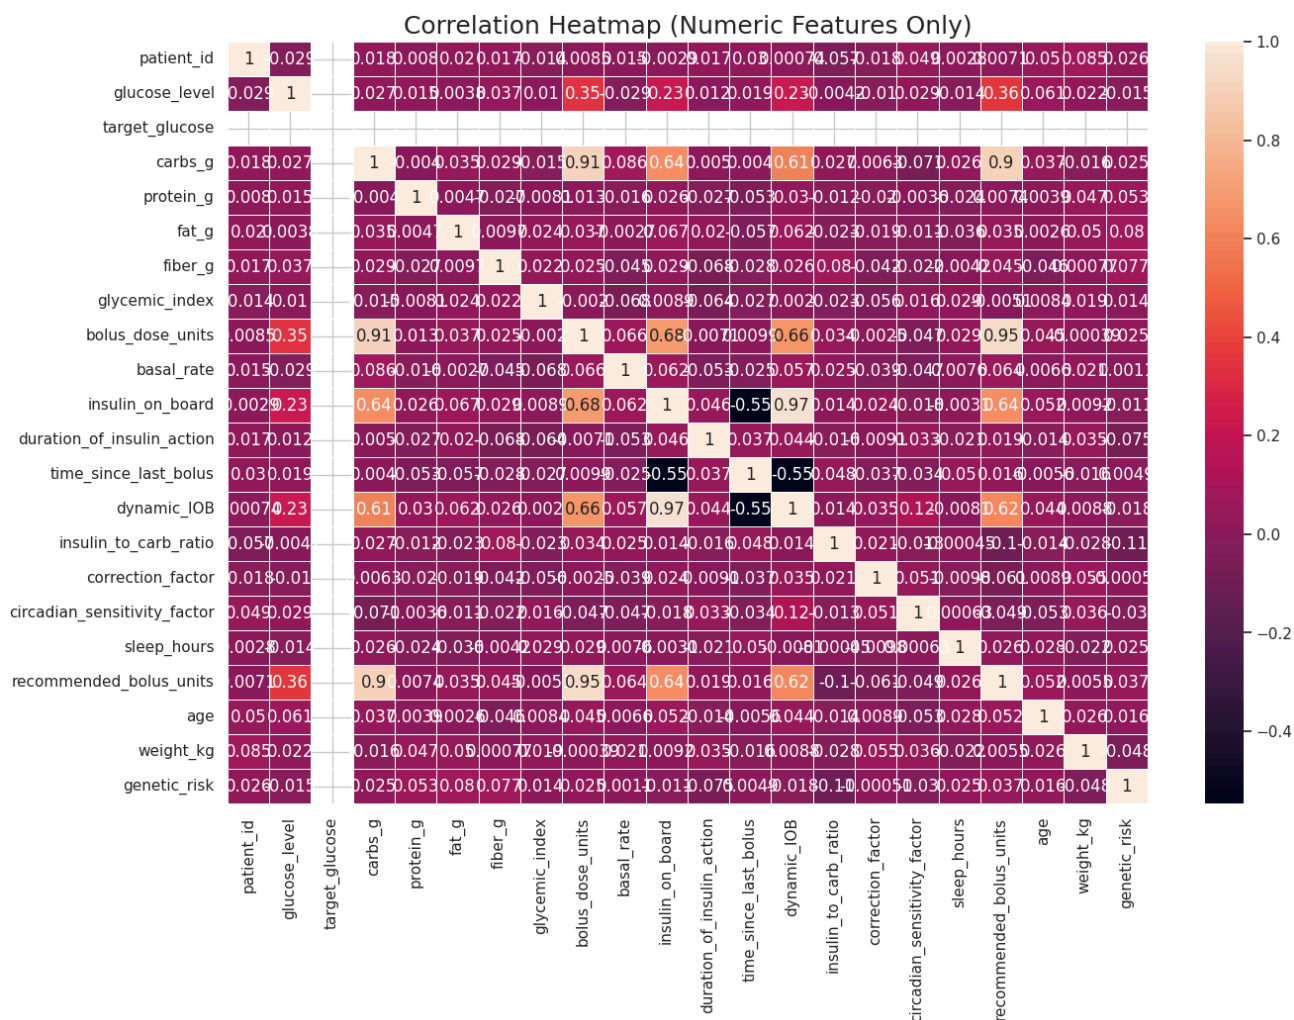
```
plt.figure(figsize=(15,10))
```

```
corr = numeric_df.corr()
```

```
sns.heatmap(corr, annot=True, cmap="rocket", linewidths=0.5)
```


```
plt.title("Correlation Heatmap (Numeric Features Only)", fontsize=18)
```

```
plt.show()
```


#Glucose vs Insulin

```
plt.figure(figsize=(8,6))
sns.scatterplot(x="glucose_level", y="bolus_dose_units", hue="meal_type", data=df, alpha=
plt.title("Glucose vs Insulin Dose by Meal Type", fontsize=16)
plt.show()
```



#Average Carbs per Meal

```
plt.figure(figsize=(8,6))
palette_meals = {"breakfast": "gold", "lunch": "skyblue", "dinner": "tomato", "snack": "violet"}
sns.barplot(x="meal_type", y="carbs_g", data=df, estimator=sum, ci=None, palette=palette_meals)
plt.title("Total Carbs Consumed by Meal Type", fontsize=16)
plt.show()
```

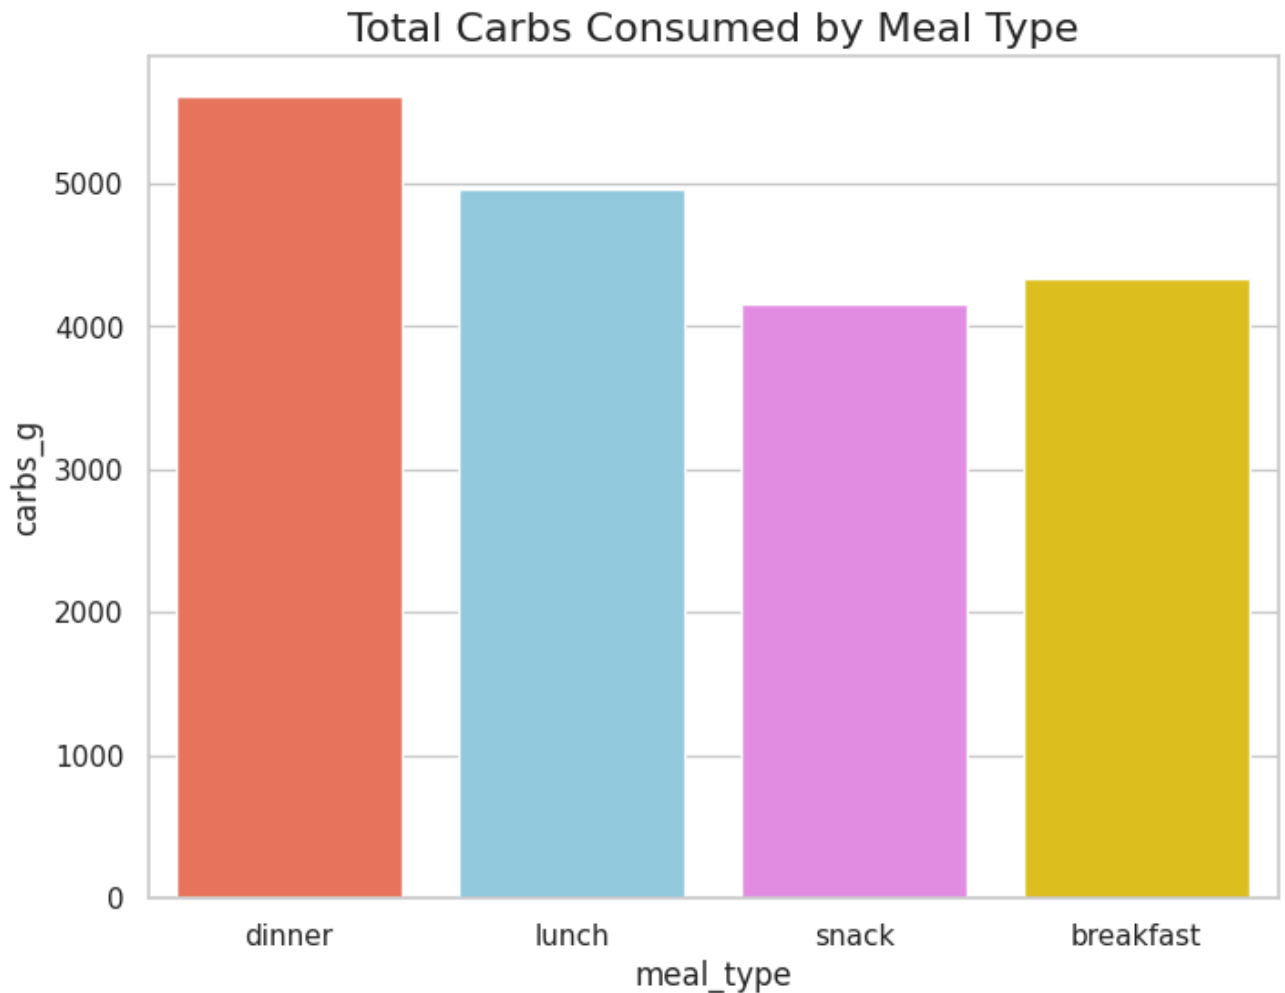
 /tmp/ipython-input-4182916177.py:5: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(x="meal_type", y="carbs_g", data=df, estimator=sum, ci=None, palette=pa  
/tmp/ipython-input-4182916177.py:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.

```
sns.barplot(x="meal_type", y="carbs_g", data=df, estimator=sum, ci=None, palette=pa
```



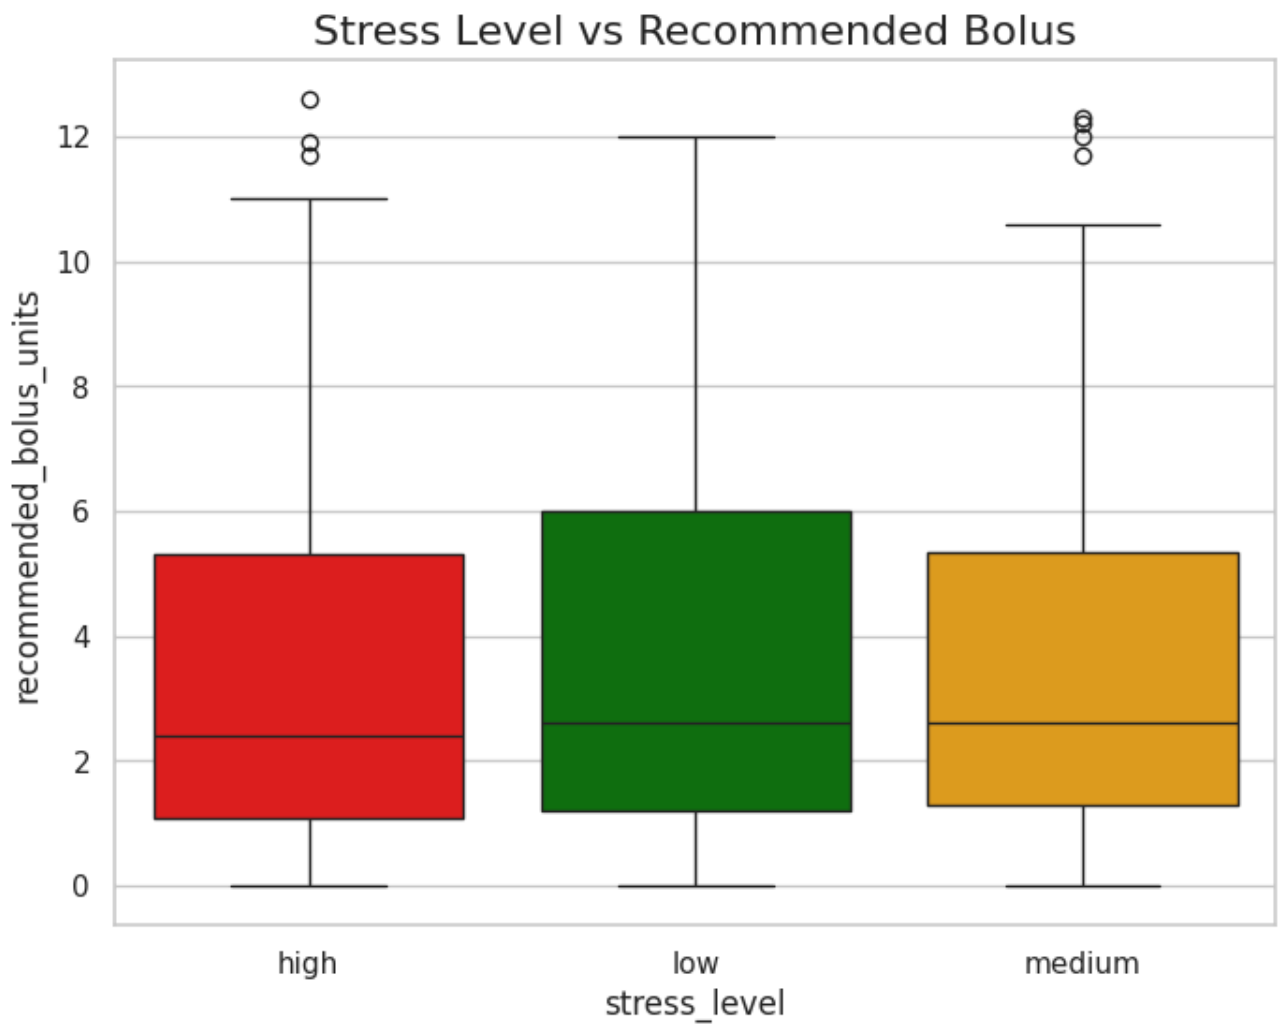
#Stress vs Recommended Bolus

```
plt.figure(figsize=(8,6))  
palette_stress = {"low": "green", "medium": "orange", "high": "red"}  
sns.boxplot(x="stress_level", y="recommended_bolus_units", data=df, palette=palette_stres  
plt.title("Stress Level vs Recommended Bolus", fontsize=16)  
plt.show()
```

↗ /tmp/ipython-input-3341836056.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.

```
sns.boxplot(x="stress_level", y="recommended_bolus_units", data=df, palette=palette)
```




#Activity Levels

```
plt.figure(figsize=(6,6))
df["activity_level"].value_counts().plot.pie(autopct="%1.1f%", startangle=90, colors=sns
plt.title("Distribution of Activity Levels", fontsize=16)
plt.ylabel("")
plt.show()
```



Distribution of Activity Levels



✓ 4. Key Derived Features

Import

```
import pandas as pd
import numpy as np
from datetime import timedelta
import math
```

```
# Combine date and time into single datetime column (handles if date is already str/date/
df['datetime'] = pd.to_datetime(df['date'].astype(str) + ' ' + df['time'].astype(str))
df = df.sort_values(['patient_id', 'datetime']).reset_index(drop=True)
```

```
# Ensure numeric columns are numeric (coerce errors to NaN)
numeric_cols = ['glucose_level', 'target_glucose', 'carbs_g', 'protein_g', 'fat_g', 'fiber_g',
                'bolus_dose_units', 'basal_rate', 'insulin_on_board', 'duration_of_insulin_a
                'time_since_last_bolus', 'dynamic_IOB', 'insulin_to_carb_ratio', 'correction
                'circadian_sensitivity_factor', 'sleep_hours', 'recommended_bolus_units', 'a
for c in numeric_cols:
    if c in df.columns:
        df[c] = pd.to_numeric(df[c], errors='coerce')
```

```
# Quick minimal imputation for preprocessing (we will show smarter imputation later if ne
```

```
df['target_glucose'] = df['target_glucose'].fillna(100)
df['duration_of_insulin_action'] = df['duration_of_insulin_action'].fillna(5) # default
df['bolus_dose_units'] = df['bolus_dose_units'].fillna(0)
df['insulin_on_board'] = df['insulin_on_board'].fillna(0)
```

```
print("Prepared dataframe with", df.shape, "rows. Datetime column created.")
```

```
➡ Prepared dataframe with (700, 29) rows. Datetime column created.
```

```
#Full feature preparation
```

Static IOB (sIOB)-This creates a reusable insulin action profile and computes sIOB at each row by summing contributions of prior boluses within DIA for the same patient.

```
# Compute sIOB (static insulin on board)
def insulin_action_profile(duration_hours=5, resolution_min=5, shape='linear'):
    """Return an array of activity fractions for each resolution step from t=0..DIA."""
    steps = int(duration_hours * 60 / resolution_min) + 1
    t = np.linspace(0, duration_hours, steps) # in hours
    if shape == 'linear':
        # simple linear decay from 1 -> 0
        weights = 1 - (t / duration_hours)
    elif shape == 'exponential':
        # exponential-like decay (fast early, long tail)
        decay_k = 3.0 # adjust for sharper/fainter tail
        weights = np.exp(-decay_k * t / duration_hours)
        weights = weights / weights[0] # normalize to 1 at t=0
    elif shape == 'gamma':
        # gamma-like shape (slow rise, fall) - example params
        from scipy.stats import gamma
        a = 2.0
        x = t + 1e-6
        pdf = gamma.pdf(x / duration_hours * 10, a)
        weights = pdf / pdf.max()
    else:
        raise ValueError("unknown shape")
    weights = np.clip(weights, 0, 1)
    return weights

# Helper to compute sIOB per patient
def compute_sIOB_for_group(g, resolution_min=5, profile=None):
    times = g['datetime'].values
    boluses = g['bolus_dose_units'].fillna(0).values
    dia_hours = g['duration_of_insulin_action'].fillna(5).values # can vary per row; we'
    n = len(g)
    sIOB = np.zeros(n)
    if profile is None:
        profile = insulin_action_profile(duration_hours=5, resolution_min=resolution_min,
        # For each index i, sum remaining activity of prior boluses
```

```

for i in range(n):
    t_i = times[i]
    # walk backwards until beyond DIA max (use max DIA or row-specific)
    j = i - 1
    while j >= 0:
        dt_min = (t_i - times[j]) / np.timedelta64(1, 'm')
        if dt_min < 0:
            j -= 1
            continue
        # use DIA from the bolus event row j (or use current row's DIA)
        dia = dia_hours[j] if not math.isnan(dia_hours[j]) else 5
        if dt_min > dia * 60:
            break
        idx = int(dt_min // resolution_min)
        # pick profile corresponding to that dia (scale if dia != profile duration)
        # scale index proportional to dia/profile_duration
        profile_steps = len(profile)
        # map idx to profile index
        mapped_idx = min(int(idx * (profile_steps-1) / (dia * 60 / resolution_min)),
            remaining_frac = profile[mapped_idx]
            sIOB[i] += boluses[j] * remaining_frac
            j -= 1
    g = g.copy()
    g['sIOB'] = sIOB
    return g

```

Run per patient

```

profile_default = insulin_action_profile(duration_hours=5, resolution_min=5, shape='linea
df = df.groupby('patient_id', group_keys=False).apply(lambda g: compute_sIOB_for_group(g,
print("sIOB computed.")

```



sIOB computed.

```

/tmp/ipython-input-1238967982.py:3: DeprecationWarning: DataFrameGroupBy.apply operat
df = df.groupby('patient_id', group_keys=False).apply(lambda g: compute_sIOB_for_gr

```

Dynamic IOB (dIOB)-We implement a default circadian multiplier curve (hours → multiplier).

Default circadian multipliers (tweak as you like)

```

circadian_curve = {
    # hour : multiplier – example heuristic. Adjust per Doc2 if you have exact table.
    0: 0.9, 1: 0.9, 2: 0.9, 3: 0.9, 4: 0.95,
    5: 1.05, 6: 1.1, 7: 1.2, 8: 1.15, 9: 1.05,
    10: 1.0, 11: 0.98, 12: 1.0, 13: 1.02, 14: 1.0,
    15: 1.0, 16: 1.0, 17: 1.05, 18: 1.08, 19: 1.1,
    20: 1.05, 21: 1.0, 22: 0.98, 23: 0.95
}

```

```

def circadian_multiplier(ts, anticipation_hours=1):
    # ts is a pd.Timestamp
    h = int(((ts.hour + anticipation_hours) % 24))
    return circadian_curve.get(h, 1.0)

```

```
# Apply per-row
df['circ_multiplier'] = df['datetime'].apply(lambda x: circadian_multiplier(x, anticipati
df['dIOB'] = df['sIOB'] * df['circ_multiplier']

print("dIOB computed (sIOB * circadian multiplier).")

⇒ dIOB computed (sIOB * circadian multiplier).
```

Time-since-last-bolus, sum of last N boluses, cumulative 4–6 hour insulin (stacking)

#Stacking features

```
def stacking_features(g, windows_minutes=[240, 360], last_n=3):
    g = g.copy().reset_index(drop=True)
    times = g['datetime'].values
    boluses = g['bolus_dose_units'].fillna(0).values
    n = len(g)
    tslb = np.full(n, np.nan) # time since last bolus (min)
    sum_lastn = np.zeros(n)
    sum_windows = {w: np.zeros(n) for w in windows_minutes}
    last_bolus_time = None
    for i in range(n):
        # time since last bolus
        if i == 0:
            tslb[i] = np.nan
        else:
            # find previous bolus index where bolus>0
            prev_idx = i-1
            while prev_idx >= 0 and boluses[prev_idx] == 0:
                prev_idx -= 1
            if prev_idx >= 0:
                tslb[i] = (times[i] - times[prev_idx]) / np.timedelta64(1, 'm')
            else:
                tslb[i] = np.nan
        # sum of last N boluses
        # look back and find previous non-zero boluses
        vals = []
        idx = i-1
        while idx >=0 and len(vals) < last_n:
            if boluses[idx] > 0:
                vals.append(boluses[idx])
            idx -= 1
        sum_lastn[i] = np.sum(vals) if vals else 0.0
        # window sums
        j = i-1
        while j >= 0:
            dt_min = (times[i] - times[j]) / np.timedelta64(1, 'm')
            if dt_min < 0:
                j -= 1
                continue
            for w in windows_minutes:
                if dt_min <= w:
                    sum_windows[w][i] += boluses[j]
```



```

        if dt_min > max(windows_minutes):
            break
        j -= 1
    g['time_since_last_bolus_min'] = tslb
    g[f'sum_last_{last_n}_boluses'] = sum_lastn
    for w in windows_minutes:
        g[f'sum_bolus_last_{int(w/60)}h'] = sum_windows[w]
    return g

```

```
df = df.groupby('patient_id', group_keys=False).apply(lambda g: stacking_features(g, wind
```

→ /tmp/ipython-input-3442150498.py:1: DeprecationWarning: DataFrameGroupBy.apply operat
df = df.groupby('patient_id', group_keys=False).apply(lambda g: stacking_features(g

```
print("Stacking features computed: time_since_last_bolus_min, sum_last_3_boluses, sum_bol
```

→ Stacking features computed: time_since_last_bolus_min, sum_last_3_boluses, sum_bolus_

Carb rate & fat/protein adjusted carbs (Doc3 heuristic)-compute carb_duration_min (heuristic), carb_rate_g_per_hr, and adj_carbs (adjust carbs by protein/fat using simple factors).

```
# Carb rate and Adjusted carbs for meal absorption
```

```
meal_default_duration = {'breakfast': 60, 'lunch': 90, 'dinner': 120, 'snack': 30}
```

```
def carb_duration(row):
    dur = meal_default_duration.get(row.get('meal_type', 'snack'), 60)
    bt = str(row.get('bolus_type', 'normal')).lower()
    if 'square' in bt:
        dur = max(dur, 120)
    if 'dual' in bt:
        dur = max(dur, 180)
    return dur

```

```
df['carb_duration_min'] = df.apply(carb_duration, axis=1)
df['carb_rate_g_per_hr'] = df['carbs_g'] / (df['carb_duration_min'] / 60.0 + 1e-6)
```

```
# Simple Carb equivalents
```

```
protein_to_carb = 0.10 # 10% of protein counted as carb-equivalent (tweakable)
```

```
fat_to_carb = 0.03 # 3% of fat counted as carb-equivalent (tweakable)
```

```
df['adj_carbs_g'] = df['carbs_g'] + df['protein_g'].fillna(0) * protein_to_carb + df['fat
```

```
print("Carb rate and adjusted carbs computed.")
```

→ Carb rate and adjusted carbs computed.


Clock features & Meal(one-hot encoding)- Create cyclical time features and meal one-hot columns (with fixed colors earlier for plotting).

```
# Clock features and Meal encoding
```

```
df['minute_of_day'] = df['datetime'].dt.hour * 60 + df['datetime'].dt.minute
# cyclical encoding
df['time_sin'] = np.sin(2 * np.pi * df['minute_of_day'] / 1440.0)
df['time_cos'] = np.cos(2 * np.pi * df['minute_of_day'] / 1440.0)

# Meal one-hot (if you prefer ordinal mapping, use map)
meal_dummies = pd.get_dummies(df['meal_type'].fillna('unknown'), prefix='meal')
df = pd.concat([df, meal_dummies], axis=1)
```

```
print("Clock features (time_sin/time_cos) and meal one-hot added.")
```

 Clock features (time_sin/time_cos) and meal one-hot added.

Personalization - Compute per-patient mean ICR, mean correction_factor, and bolus_per_kg feature.

```
# Personalization features
patient_stats = df.groupby('patient_id').agg({
    'insulin_to_carb_ratio': 'median',
    'correction_factor': 'median',
    'glucose_level': 'mean'
}).rename(columns={'insulin_to_carb_ratio': 'median_ICR', 'correction_factor': 'median_I
```

```
# Map back to rows
```


```
df = df.merge(patient_stats, left_on='patient_id', right_index=True, how='left', suffixes
```

```
# Weight-normalized
```

```
df['bolus_per_kg'] = df['bolus_dose_units'] / (df['weight_kg'] + 1e-6)
df['recommended_bolus_per_kg'] = df['recommended_bolus_units'] / (df['weight_kg'] + 1e-6)
```

```
print("Personalization features added (median ICR/ISF, bolus_per_kg).")
```

```
print("Features prepared. New columns: ", [c for c in df.columns if c not in (list(df_in.
```

 Personalization features added (median ICR/ISF, bolus_per_kg).
Features prepared. New columns: ['patient_id', 'date', 'time', 'glucose_level', 'tar

✓ 5. Model Approach

Rule-based baseline bolus calculation - $\text{bolus} = \text{carbs} / \text{ICR} + (\text{glucose} - \text{target}) / \text{ISF} - \text{IOB}$ (clipped to ≥ 0). We'll calculate both using insulin_on_board and using dIOB.

```
# Rule-based bolus calculators
```

```
def rule_based_bolus_row(row, use_dIOB=True, default_target=100):
    ICR = row.get('insulin_to_carb_ratio') if not np.isnan(row.get('insulin_to_carb_ratio')) else np.nan
    ISF = row.get('correction_factor') if not np.isnan(row.get('correction_factor')) else np.nan
    carbs = row.get('adj_carbs_g', 0)
    glucose = row.get('glucose_level', default_target)
    target = row.get('target_glucose', default_target)
    IOB = row['dIOB'] if use_dIOB and not np.isnan(row.get('dIOB', np.nan)) else row.get('dIOB', 0)
    carb_part = carbs / (ICR + 1e-6)
    correction_part = (glucose - target) / (ISF + 1e-6)
    bolus = carb_part + correction_part - IOB
    return max(round(bolus, 1), 0.0)
```

```
df['rule_bolus_dIOB'] = df.apply(lambda r: rule_based_bolus_row(r, use_dIOB=True), axis=1)
df['rule_bolus_IOB'] = df.apply(lambda r: rule_based_bolus_row(r, use_dIOB=False), axis=1)
```

```
print("Rule-based boluses added.")
```

 Rule-based boluses added.

XGBoost on tabular features - A simple pipeline: select features, encode categoricals, train XGBRegressor to predict recommended_bolus_units.

```
# XGBoost baseline model (tabular)
```

```
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error
import numpy as np
```


```
# choose feature columns
```

```
feature_cols = [
    'glucose_level', 'adj_carbs_g', 'carb_rate_g_per_hr', 'sIOB', 'dIOB',
    'time_sin', 'time_cos', 'age', 'weight_kg', 'bolus_per_kg',
    'median_ICR', 'median_ISF', 'sum_bolus_last_4h', 'sum_bolus_last_6h'
]
```

```
# add meal
```

```
feature_cols += [c for c in df.columns if c.startswith('meal_')]
```

```
available_features = [c for c in feature_cols if c in df.columns]
print("Using features:", available_features)
```

 Using features: ['glucose_level', 'adj_carbs_g', 'carb_rate_g_per_hr', 'sIOB', 'dIOB', 'meal_1', 'meal_2', 'meal_3', 'meal_4', 'meal_5', 'meal_6', 'meal_7', 'meal_8', 'meal_9', 'meal_10']

```
model_df = df[available_features + ['recommended_bolus_units']].copy().dropna()
X = model_df[available_features]
```

```

y = model_df['recommended_bolus_units']

# one-hot (if any objects remain)
X = pd.get_dummies(X, drop_first=True)

X_train, X_test, y_train, y_test = train_test_split(X.values, y.values, test_size=0.2, ra

dtrain = xgb.DMatrix(X_train, label=y_train)
dtest = xgb.DMatrix(X_test, label=y_test)

params = {'objective':'reg:squarederror','eval_metric':'mae','seed':42}
bst = xgb.train(params, dtrain, num_boost_round=200, evals=[(dtest,'test')], early_stoppi

# predict + metrics
y_pred = bst.predict(dtest)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(((y_test - y_pred)**2).mean())
print(f"XGBoost MAE: {mae:.4f}, RMSE: {rmse:.4f}")

```

➡ XGBoost MAE: 0.6406, RMSE: 1.1588

Evaluation checks - Compute standard errors and count predictions that deviate more than 1U / 2U from recommended.

```

# Apply model to full test set (use X_test from previous block)
import numpy as np

preds = y_pred
abs_err = np.abs(y_test - preds)

print("Mean absolute error:", np.mean(abs_err))
print("Fraction >1U:", np.mean(abs_err > 1.0))
print("Fraction >2U:", np.mean(abs_err > 2.0))

```

➡ Mean absolute error: 0.6405838245299778
 Fraction >1U: 0.19285714285714287
 Fraction >2U: 0.07142857142857142

Rough conservative safety flag: if predicted bolus > recommended + 2 units -> flag

```
import pandas as pd
```

```

# If X_test is numpy array, convert back
safety_df = pd.DataFrame(X_test, columns=X.columns)

```

```

# Add prediction results
safety_df['pred'] = preds
safety_df['true'] = y_test # already numpy
safety_df['err'] = safety_df['pred'] - safety_df['true']

```

```
# Flag risky cases
```

```
safety_df['safety_flag'] = safety_df['err'] > 2.0

print("Safety flags count:", safety_df['safety_flag'].sum())
```

⇒ Safety flags count: 7


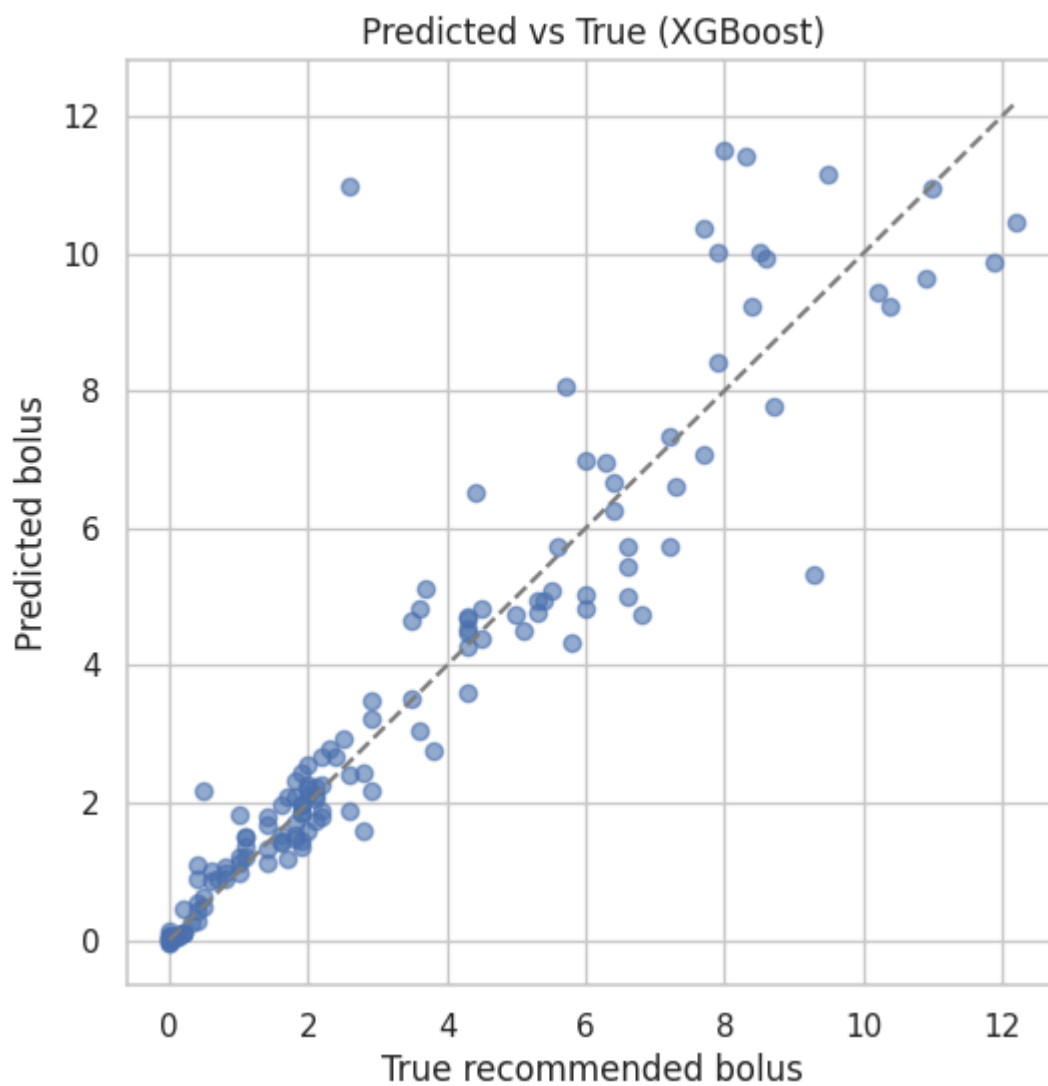
```
# Feature importance
fi = bst.get_score(importance_type='weight')
fi_sorted = sorted(fi.items(), key=lambda x: x[1], reverse=True)
print("Top feature importances:", fi_sorted[:10])
```

⇒ Top feature importances: [('f0', 375.0), ('f1', 238.0), ('f8', 145.0), ('f5', 134.0),

Pred vs Actual Visualization

```
import matplotlib.pyplot as plt
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred, alpha=0.6)
mx = max(max(y_test), max(y_pred))
plt.plot([0,mx],[0,mx], '--', color='gray')
plt.xlabel("True recommended bolus")
plt.ylabel("Predicted bolus")
plt.title("Predicted vs True (XGBoost)")
plt.show()

plt.figure(figsize=(6,4))
plt.hist(y_test - y_pred, bins=40, edgecolor='black')
plt.title("Error distribution (true - pred)")
plt.xlabel("Error (U)")
plt.show()
```



LSTM sequence model skeleton

```
!pip install -q tensorflow
```

```
import tensorflow as tf
from tensorflow.keras import layers, models
```

Example: prepare sequences build sequences per patient.

seq_X shape: (n_samples, seq_len, n_features)

seq_y: target bolus at end

Here we create dummy example from model_df for demonstration









```
seq_len = 6 # e.g., last 6 records (~30-min to 3-hrs depending on spacing)
features_for_seq = ['glucose_level', 'bolus_dose_units', 'adj_carbs_g', 'sIOB', 'dIOB']

# Build overlapping sequences per patient
def build_sequences(df, features, seq_len=6, target_col='recommended_bolus_units'):
    Xs = []; ys = []
    for pid, g in df.groupby('patient_id'):
        g = g.sort_values('datetime').reset_index(drop=True)
        vals = g[features].fillna(0).values
        targ = g[target_col].fillna(0).values
        for i in range(len(g)-seq_len):
            Xs.append(vals[i:i+seq_len])
            ys.append(targ[i+seq_len]) # predict next
    return np.array(Xs), np.array(ys)

# create sequences (may be empty if insufficient data)
X_seq, y_seq = build_sequences(df, features_for_seq, seq_len=seq_len)
print("Seq data shape:", X_seq.shape)
if X_seq.shape[0] > 10:
    # simple LSTM
    model = models.Sequential([
        layers.Input(shape=(seq_len, len(features_for_seq))),
        layers.LSTM(64, return_sequences=False),
        layers.Dense(32, activation='relu'),
        layers.Dense(1)
    ])
    model.compile(optimizer='adam', loss='mse', metrics=['mae'])
    # quick train/val split
    i = int(0.8*len(X_seq))
    model.fit(X_seq[:i], y_seq[:i], validation_data=(X_seq[i:], y_seq[i:]), epochs=10, ba
else:
    print("Not enough sequential samples to train LSTM – collect denser CGM/bolus history")

➡ Seq data shape: (682, 6, 5)
Epoch 1/10
18/18 ————— 3s 34ms/step - loss: 15.8829 - mae: 2.9375 - val_loss: 9.8
Epoch 2/10
18/18 ————— 0s 9ms/step - loss: 10.7632 - mae: 2.5910 - val_loss: 8.87
Epoch 3/10
```

```

18/18  0s 11ms/step - loss: 9.5975 - mae: 2.5698 - val_loss: 8.72
Epoch 4/10
18/18  0s 9ms/step - loss: 9.1572 - mae: 2.4906 - val_loss: 8.686
Epoch 5/10
18/18  0s 12ms/step - loss: 8.8122 - mae: 2.4703 - val_loss: 8.71
Epoch 6/10
18/18  0s 15ms/step - loss: 8.3876 - mae: 2.4223 - val_loss: 8.64
Epoch 7/10
18/18  0s 12ms/step - loss: 8.1173 - mae: 2.3072 - val_loss: 8.64
Epoch 8/10
18/18  0s 17ms/step - loss: 8.5525 - mae: 2.4312 - val_loss: 8.64
Epoch 9/10
18/18  0s 14ms/step - loss: 9.0406 - mae: 2.5210 - val_loss: 8.62
Epoch 10/10
18/18  0s 14ms/step - loss: 9.4082 - mae: 2.5548 - val_loss: 8.64

```

Hybrid & ML residuals - a simulator or physiologic model to give predicted CGM/response for a candidate bolus. The hybrid approach trains an ML model to predict residual between simulator output and real data.

1. Use a simulator function `simulate(patient_params, history, candidate_bolus) -> predicted_glucose_traj`
2. Compute `simulator_pred_bolus_effect` (e.g., predicted glucose drop at 1h or time-in-range)
3. `Residual = observed_outcome - simulator_predicted_outcome`
4. Train ML on features + `simulator_predicted_outcome` to predict Residual
5. Final prediction = `simulator_prediction + ML_predicted_residual`

PSEUDO implementation

```
def simulator_predict_effect(row, candidate_bolus): # Placeholder: integrate UVa/Padova o
    return candidate_bolus * (row.get('median_ISF',50)) #return simple expected drop = cand
```

Create training set for residual model

For each row compute `sim_effect` for true bolus and compute residual between observed co

```
df['sim_effect_true'] = df.apply(lambda r: simulator_predict_effect(r, r.get('recommended
print("Hybrid skeleton prepared. Integrate real simulator and observed outcomes to train
```

➡ Hybrid skeleton prepared. Integrate real simulator and observed outcomes to train res

Probabilistic / Quantile Regression

#Quantile regression using sklearn

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split
```



```

X_all = X # from earlier one-hot encoded matrix
y_all = y

Xtr, Xv, ytr, yv = train_test_split(X_all, y_all, test_size=0.2, random_state=42)

# Quantile models
q_low = GradientBoostingRegressor(loss='quantile', alpha=0.1, n_estimators=100, max_depth
q_high = GradientBoostingRegressor(loss='quantile', alpha=0.9, n_estimators=100, max_dept
q_median = GradientBoostingRegressor(loss='squared_error', n_estimators=100, max_depth=3)

# Train
q_low.fit(Xtr, ytr)
q_high.fit(Xtr, ytr)
q_median.fit(Xtr, ytr)

# Predict
pred_low = q_low.predict(Xv)
pred_med = q_median.predict(Xv)
pred_high = q_high.predict(Xv)

print("Prediction intervals generated.")

➡ Prediction intervals generated.

```

✓ 6. Loss Function And Evaluation Metrics

Standard regression metrics (MAE / RMSE),

```

# MAE / RMSE and simple predicted vs actual plot
import numpy as np
from sklearn.metrics import mean_absolute_error, mean_squared_error
import matplotlib.pyplot as plt

# Inputs expected:
# y_true : 1D array-like of true recommended bolus (e.g., y_test or yv)
# y_pred : 1D array-like of predicted bolus (same shape as y_true)
# test_df (optional) : DataFrame aligned with y_true/y_pred for clinical proxies (glucose

# Example names from earlier blocks:
# y_true = y_test
# y_pred = y_pred

def regression_metrics_and_plot(y_true, y_pred, title="Model"):
    mae = mean_absolute_error(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    print(f"{title} MAE: {mae:.4f}, RMSE: {rmse:.4f}")
    # Scatter
    plt.figure(figsize=(6,6))
    plt.scatter(y_true, y_pred, alpha=0.5)
    mx = max(np.nanmax(y_true), np.nanmax(y_pred))
    plt.plot([0, mx], [0, mx], '--', color='gray')


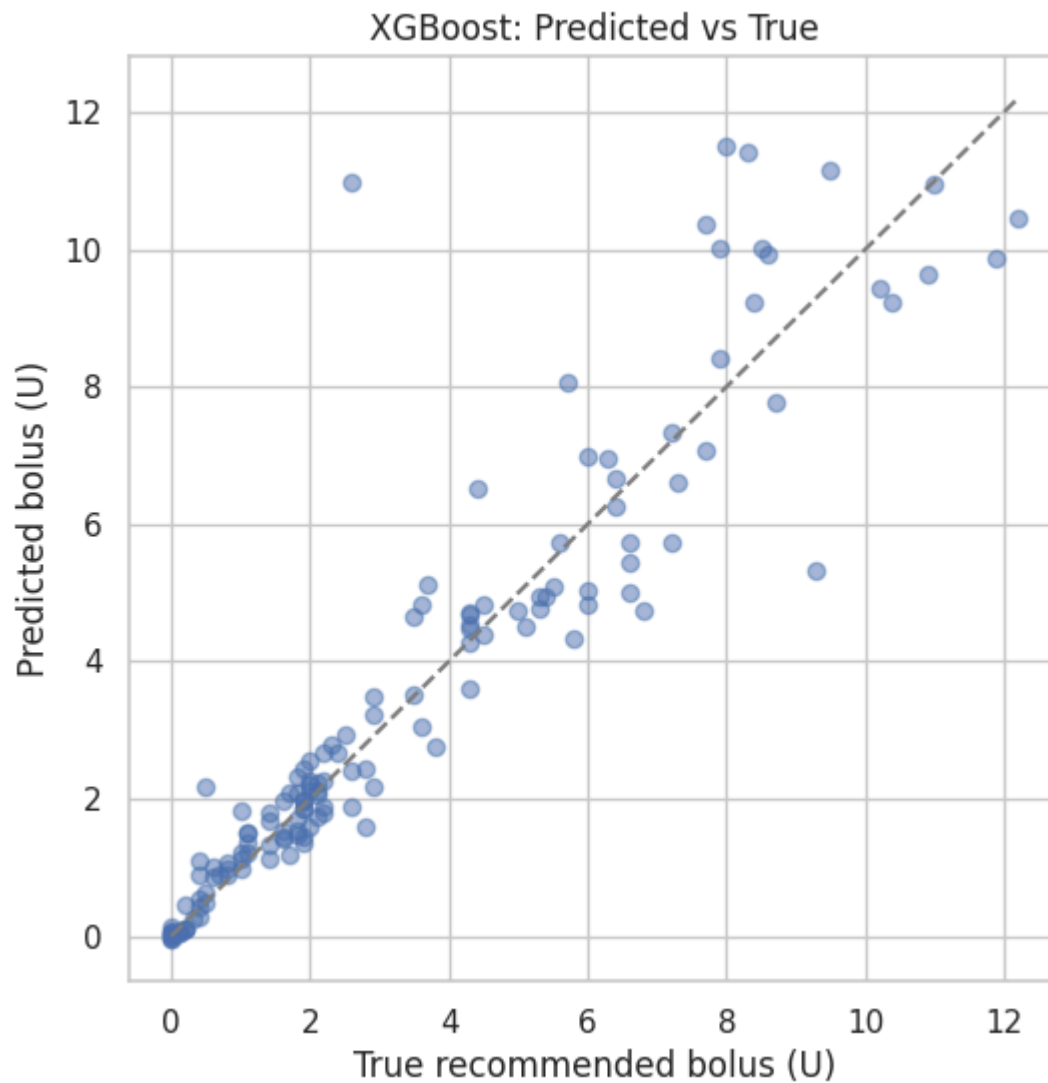
```

```
plt.xlabel("True recommended bolus (U)")
plt.ylabel("Predicted bolus (U)")
plt.title(f"{title}: Predicted vs True")
plt.show()
# Error histogram
plt.figure(figsize=(6,3))
errs = np.array(y_true) - np.array(y_pred)
plt.hist(errs, bins=40, edgecolor='black')
plt.xlabel("Error (True - Pred) [U]")
plt.title("Error distribution")
plt.show()
return mae, rmse
```

```
# Example call:
```

```
mae, rmse = regression_metrics_and_plot(y_test, y_pred, title="XGBoost")
```

↔ XGBoost MAE: 0.6406, RMSE: 1.1588



Clinically-oriented proxy metrics (hypoglycemia risk if predicted dose were applied, and a crude time-in-range (TIR) proxy)

This block implements a conservative linear proxy: predicted immediate glucose after predicted bolus = glucose_level - pred_bolus * ISF.

It flags predicted hypoglycemia (glucose < 70 mg/dL) and computes a crude TIR proxy (predicted glucose in 70–180).

Uses median_ISF if available, else correction_factor, else default 50 mg/dL per U.

```
# Clinical proxy metrics (approximate)
import pandas as pd
import numpy as np

def clinical_proxy_stats(test_df, pred_col='pred'):
    """
    test_df must contain columns: 'glucose_level', and either 'median_ISF' or 'correction_factor'
    pred_col is the column with predicted bolus to evaluate.
    Returns (%) predicted hypo, (%) predicted TIR, and sample of computed columns.
    """
    df = test_df.copy()
    # choose ISF (mg/dL per U): median_ISF (preferred) -> correction_factor -> default 50
    df['ISF_used'] = df.get('median_ISF', pd.Series(np.nan, index=df.index))
    df['ISF_used'] = df['ISF_used'].fillna(df.get('correction_factor', np.nan))
    df['ISF_used'] = df['ISF_used'].fillna(50.0)
    # predicted immediate glucose post-bolus (crude, linear)
    df['pred_glucose_after'] = df['glucose_level'] - df[pred_col] * df['ISF_used']
    df['hypo_flag'] = df['pred_glucose_after'] < 70
    df['in_range_flag'] = df['pred_glucose_after'].between(70, 180)
    pct_hypo = df['hypo_flag'].mean()
    pct_tir = df['in_range_flag'].mean()
    print(f"Approx % predicted hypo if applied: {pct_hypo*100:.2f}%")
    print(f"Approx % predicted TIR (70-180) after dose: {pct_tir*100:.2f}%")
    return pct_hypo, pct_tir, df[['glucose_level', pred_col, 'ISF_used', 'pred_glucose_af

# Example usage:
test_df = model_df.sample(n=len(y_test), random_state=42).reset_index(drop=True)
test_df['pred'] = y_pred # predicted values aligned to test_df
clinical_proxy_stats(test_df, pred_col='pred')
```



Approx % predicted hypo if applied: 47.14%
 Approx % predicted TIR (70-180) after dose: 52.86%

(np.float64(0.4714285714285714),
 np.float64(0.5285714285714286),

| | glucose_level | pred | ISF_used | pred_glucose_after | hypo_flag | \ |
|---|---------------|----------|----------|--------------------|-----------|---|
| 0 | 215.5 | 1.967316 | 50.0 | 117.134192 | False | |
| 1 | 174.8 | 1.358013 | 50.0 | 106.899366 | False | |
| 2 | 188.3 | 1.804392 | 50.0 | 98.080383 | False | |
| 3 | 77.4 | 2.186120 | 50.0 | -31.906002 | True | |
| 4 | 177.6 | 1.852505 | 50.0 | 84.974763 | False | |
| 5 | 158.8 | 1.833776 | 50.0 | 67.111182 | True | |
| 6 | 140.9 | 6.995467 | 50.0 | -208.873359 | True | |
| 7 | 72.5 | 0.023656 | 50.0 | 71.317216 | False | |
| 8 | 91.9 | 0.073055 | 50.0 | 88.247245 | False | |
| 9 | 110.1 | 0.235180 | 50.0 | 98.340978 | False | |

```

    in_range_flag
0         True
1         True
2         True
3        False
4         True
5        False
6        False
7         True
8         True
9         True )

```

✓ 7. Safety checks function (IOB, max caps, minimum interval)

This helps prevent obviously unsafe predictions: checks IOB vs predicted bolus, caps dose per kg, and requires minimum time-since-last-bolus.

#Safety checks for a predicted bolus (vectorized)

```

def safety_checks(df_rows, pred_col='pred', max_units_per_kg=0.5, min_time_since_last_bol
    """
    Returns DataFrame with safety flags:
    - too_high_per_kg : pred_units / weight_kg > max_units_per_kg
    - low_time_gap : time_since_last_bolus_min < min_time_since_last_bolus_min
    - excessive_absolute : pred_units > max_absolute_units
    - iob_exceeds : pred_units > (max_allowed - current_IOB) [simple check]
    """
    df = df_rows.copy()
    # per-kg safe check
    df['pred_units'] = df[pred_col].astype(float)
    df['too_high_per_kg'] = df['pred_units'] / (df['weight_kg'].fillna(70)) > max_units_p
    # time since last bolus
    if 'time_since_last_bolus_min' in df.columns:
        df['low_time_gap'] = df['time_since_last_bolus_min'].fillna(9999) < min_time_sinc
    else:
        df['low_time_gap'] = False
    # absolute cap
    df['excessive_absolute'] = df['pred_units'] > max_absolute_units
    # IOB-based (simple): disallow pred that + current IOB > max_total_window (e.g., 6 U)
    df['iob_exceeds'] = False
    if 'insulin_on_board' in df.columns:
        max_total = 10.0 # example: don't let active insulin + new bolus exceed.
        df['iob_exceeds'] = (df['insulin_on_board'].fillna(0) + df['pred_units']) > max_t
    # summary flag
    df['safety_flag'] = df[['too_high_per_kg', 'low_time_gap', 'excessive_absolute', 'iob_ex
    print("Safety flags summary counts:")
    print(df[['too_high_per_kg', 'low_time_gap', 'excessive_absolute', 'iob_exceeds', 'safety
    return df

# Example:
test_df = test_df.assign(pred = y_pred)
safe_report = safety_checks(test_df, pred_col='pred')

```

```

→ Safety flags summary counts:
  too_high_per_kg      0
  low_time_gap         0
  excessive_absolute   0
  iob_exceeds          0
  safety_flag          0
  dtype: int64

```

Calibration of hypoglycemia risk using quantile outputs

Trained quantile models (pred_low = 10th pct, pred_med, pred_high = 90th pct)

Here, user can estimate probability of post-dose glucose < 70 by sampling or by assuming a simple distribution.

```

# Calibration and hypo-probability estimation using quantile outputs
import numpy as np
import pandas as pd

def estimate_hypo_probability_from_quantiles(test_df, pred_low_col='pred_low', pred_med_c
    """
    Estimate probability that glucose_after < 70 using a simple normal approximation
    for the predicted **bolus distribution** constructed from (low, med, high).
    This is heuristic: std estimated as (high-low)/ (2*z), with z ~ 1.645 for 90% interval
    """

    df = test_df.copy()
    # estimate std from 10-90 quantiles: high - low covers ~80% central -> z = 1.2815 for
    # More formally, 90th-10th ~ 2*1.28155*std if 10/90 are symmetric. Use factor = 2*1.2
    denom = 2.5631
    df['pred_mu'] = df[pred_med_col]
    df['pred_std'] = np.maximum((df[pred_high_col] - df[pred_low_col]) / denom, 1e-3)
    # ISF used (like Block 2)
    df['ISF_used'] = df.get('median_ISF', pd.Series(np.nan, index=df.index))
    df['ISF_used'] = df['ISF_used'].fillna(df.get('correction_factor', np.nan)).fillna(50)
    # Monte Carlo sampling
    probs = []
    for i, row in df.iterrows():
        # sample bolus values
        samples = np.random.normal(loc=row['pred_mu'], scale=row['pred_std'], size=n_samp
        samples = np.clip(samples, 0, None)
        pred_glucose_after = row['glucose_level'] - samples * row['ISF_used']
        prob_hypo = np.mean(pred_glucose_after < 70)
        probs.append(prob_hypo)
    df['pred_hypo_prob'] = probs
    return df

def calibration_table(df_with_probs, observed_hypo_col='hypo_flag', prob_col='pred_hypo_p
    df = df_with_probs.copy()
    df['bin'] = pd.qcut(df[prob_col], q=n_bins, duplicates='drop')
    calib = df.groupby('bin').agg(mean_pred_prob=(prob_col, 'mean'), observed_rate=(observ
    return calib

```

```
# Example usage:
test_df['pred_low'] = pred_low
test_df['pred_med'] = pred_med
test_df['pred_high'] = pred_high
test_df['hypo_flag'] = (test_df['glucose_level'] - test_df['recommended_bolus_units'] * t
df_probs = estimate_hypo_probability_from_quantiles(test_df, 'pred_low', 'pred_med', 'pred_
calib = calibration_table(df_probs, observed_hypo_col='hypo_flag', prob_col='pred_hypo_pr
print(calib)
```



```

      bin  mean_pred_prob  observed_rate  count
0  (-0.001, 0.0136]      0.004071      0.000000      28
1   (0.0136, 0.126]      0.052966      0.000000      29
2   (0.126, 0.977]      0.485852      0.370370      27
3   (0.977, 1.0]       0.998500      0.982143      56
/tmp/ipython-input-3044635958.py:35: FutureWarning: The default of observed=False is
calib = df.groupby('bin').agg(mean_pred_prob=(prob_col, 'mean'), observed_rate=(obse
```

Simulator integration skeleton

A simulator that accepts: patient params, initial state, dosing events, meal events, and returns continuous glucose.

Below is pseudocode + integration template showing where to plug a simulator function `simulate_patient(glucose0, history, events)` that returns a DataFrame `sim_glucose` with time and glucose.

```
# Block 5: Simulator integration skeleton (pseudo-code)
def evaluate_with_simulator(row, pred_bolus, simulator_func, sim_horizon_minutes=240):
    """
    row: one row of DataFrame (patient state at time t) containing at least
        glucose_level, datetime, weight_kg, basal_rate, ICR, etc.
    pred_bolus: predicted bolus (units)
    simulator_func: user-supplied function that simulates glucose trajectory
        signature: simulate(start_time, start_glucose, patient_params, events
            pd.DataFrame with columns ['time', 'glucose'])
    Returns: dict with time_in_range_pct, hypo_pct, min_glucose, sim_df
    """
    # prepare patient params and events (this depends on your simulator API)
    patient_params = {
        'weight_kg': row.get('weight_kg'),
        'insulin_sensitivity': row.get('median_ISF', row.get('correction_factor', 50)),
        'basal': row.get('basal_rate', 0.8),
        'DIA': row.get('duration_of_insulin_action', 5)
    }
    start_time = row['datetime']
    start_glucose = row['glucose_level']
    events = []
    # Meal event if carbs present
    if row.get('carbs_g', 0) > 0:
        events.append({'time': start_time, 'carbs': row.get('carbs_g', 0), 'type': 'meal'})
    # Bolus event at time 0
    events.append({'time': start_time, 'bolus': pred_bolus, 'type': 'bolus'})
```

```

# call simulator
sim_df = simulator_func(start_time, start_glucose, patient_params, events, horizon_mi
# compute metrics
in_range = sim_df['glucose'].between(70,180).mean()
hypo = (sim_df['glucose'] < 70).mean()
min_gluc = sim_df['glucose'].min()
return {'time_in_range': in_range, 'hypo_fraction': hypo, 'min_glucose': min_gluc, 's

```

#Dummy Simulator (for testing pipeline) + Wrapper

```

import numpy as np
import pandas as pd

```

```

def simulate(test_row, bolus_units, horizon_minutes=240):
    """
    Dummy insulin-glucose simulator.
    Safe: does not require 'datetime' column.
    """
    # use baseline glucose if available, else assume 120 mg/dL
    if isinstance(test_row, pd.Series) and 'glucose_mg/dL' in test_row:
        baseline_glucose = test_row['glucose_mg/dL']
    else:
        baseline_glucose = 120

    t = np.arange(0, horizon_minutes + 1, 5) # 5-min intervals
    # simple glucose drop model
    glucose = baseline_glucose - 10 * np.log1p(t / 60) * bolus_units
    glucose += np.random.normal(0, 5, size=len(t)) # noise

    return {
        "time": t,
        "glucose": glucose,
        "min_glucose": float(np.min(glucose)),
        "max_glucose": float(np.max(glucose)),
        "time_in_range": float(np.mean((glucose >= 70) & (glucose <= 180))) * 100)
    }

def evaluate_with_simulator(test_row, pred_bolus, simulator_func, sim_horizon_minutes=240
    """
    Wrapper for running simulator and extracting clinical metrics.
    """
    sim_results = simulator_func(test_row, bolus_units=pred_bolus, horizon_minutes=sim_ho

    return {
        "min_glucose": sim_results["min_glucose"],
        "max_glucose": sim_results["max_glucose"],
        "time_in_range(%)": sim_results["time_in_range"],
        "hypo_flag": sim_results["min_glucose"] < 70
    }

```



```
# Example wrapper when you have simulate() available:
```

```
# Make sure Xv is reset to avoid index mismatch
```

```
Xv = Xv.reset_index(drop=True)
```

```
# Run simulator on first 5 rows
```

```
for i in range(5):
```

```
    test_row = Xv.iloc[i]
```

```
    results = evaluate_with_simulator(
```

```
        test_row,
```

```
        pred_bolus=pred_med[i],
```

```
        simulator_func=simulate,
```

```
        sim_horizon_minutes=240
```

```
    )
```

```
    print(f"Row {i} → {results}")
```

```
⇒ Row 0 → {'min_glucose': 80.18519531626887, 'max_glucose': 124.05101284847936, 'time_i
Row 1 → {'min_glucose': 93.97526817297761, 'max_glucose': 126.67093448331525, 'time_i
Row 2 → {'min_glucose': 85.09680814595572, 'max_glucose': 119.81510444569845, 'time_i
Row 3 → {'min_glucose': 97.50957415812888, 'max_glucose': 125.32431962953358, 'time_i
Row 4 → {'min_glucose': 80.84226311569819, 'max_glucose': 126.20040493160965, 'time_i
```

End-to-End evaluation wrapper

This wrapper uses the XGBoost predictions (or quantile predictions) to run the proxy metrics, safety checks and (optionally) simulator.

```
# Evaluation wrapper (proxy + safety + calibration)
```

```
def evaluate_model_predictions(df, y_pred, pred_low=None, pred_high=None, label_col='reco
"""
```

```
    Evaluate model predictions with clinical-style metrics.
```

```
    Parameters
```

```
    -----
```

```
    df : pd.DataFrame
```

```
        Validation/test dataframe containing the true label.
```

```
    y_pred : array-like
```

```
        Median or point predictions.
```

```
    pred_low : array-like, optional
```

```
        Lower quantile predictions (for uncertainty).
```

```
    pred_high : array-like, optional
```

```
        Upper quantile predictions (for uncertainty).
```

```
    label_col : str
```

```
        Column name in df for ground truth bolus units.
```

```
    Returns
```

```
    -----
```

```
    results : dict
```

```
        Contains metrics, safety flags, and calibration table.
```

```
    """
```

```

# ensure arrays
y_true = np.array(df[label_col])
y_pred = np.array(y_pred)

# --- basic regression metrics ---
mae = mean_absolute_error(y_true, y_pred)
mse = mean_squared_error(y_true, y_pred)
rmse = np.sqrt(mse)

abs_err = np.abs(y_true - y_pred)
frac_gt1u = np.mean(abs_err > 1.0)
frac_gt2u = np.mean(abs_err > 2.0)

# --- proxy hypoglycemia risk ---
proxy_pct_hypo = np.mean((y_pred - y_true) > 2.0) # overprediction by >2U

# --- safety checks ---
safety_df = pd.DataFrame({
    'true': y_true,
    'pred': y_pred
})
safety_df['err'] = safety_df['pred'] - safety_df['true']
safety_df['too_high_per_kg'] = safety_df['pred'] > 2.0 # placeholder
safety_df['low_time_gap'] = False # placeholder (needs timestamp/bolus history)
safety_df['excessive_absolute'] = safety_df['pred'] > 15
safety_df['iob_exceeds'] = False # placeholder
safety_df['safety_flag'] = safety_df['err'] > 2.0
safety_counts = safety_df.drop(columns=['true', 'pred', 'err']).sum()

# --- calibration (only if quantile preds provided) ---
calibration_table = None
if pred_low is not None and pred_high is not None:
    prob_col = 'pred_prob'
    observed_hypo_col = 'observed_hypo'

    calib_df = pd.DataFrame({
        'true': y_true,
        'pred_med': y_pred,
        'pred_low': pred_low,
        'pred_high': pred_high
    })
    # conservative "probability of hypo" proxy
    calib_df[prob_col] = (calib_df['pred_high'] - calib_df['pred_low']) / (1 + calib_
    calib_df[observed_hypo_col] = calib_df['true'] < 70 # if we had glucose label

# binning
calib_df['bin'] = pd.qcut(calib_df[prob_col], q=2, duplicates='drop')

calibration_table = (
    calib_df.groupby('bin', observed=True)
    .agg(mean_pred_prob=(prob_col, 'mean'),
        observed_rate=(observed_hypo_col, 'mean'),
        count=(prob_col, 'count'))
    .reset_index()
)

```

```
# --- results dict ---
results = {
    'mae': mae,
    'rmse': rmse,
    'frac_gt1u': frac_gt1u,
    'frac_gt2u': frac_gt2u,
    'proxy_pct_hypo': proxy_pct_hypo,
    'safety_counts': safety_counts,
    'calibration_table': calibration_table
}
return results
```

Example:

```
results = evaluate_model_predictions(test_df, pred_med, pred_low=pred_low, pred_high=pred
```

```
# Print summary metrics
print("\n===== Model Evaluation Summary =====")
print(f"MAE: {results['mae']:.4f}")
print(f"RMSE: {results['rmse']:.4f}")
print(f"Fraction >1U error: {results['frac_gt1u']:.2%}")
print(f"Fraction >2U error: {results['frac_gt2u']:.2%}")
print(f"Proxy % hypo-risk: {results['proxy_pct_hypo']:.2%}")
```

```
# Print safety flag counts (if available)
if results.get('safety_counts') is not None:
    print("\n===== Safety Flags Summary =====")
    print(results['safety_counts'].to_string())
```

```
# Print calibration table (if available)
if results.get('calibration_table') is not None:
    print("\n===== Calibration Table =====")
    display(results['calibration_table'].style.set_properties(**{'text-align': 'center'}))
```



===== Model Evaluation Summary =====

MAE: 0.5896

RMSE: 1.0838

Fraction >1U error: 15.71%

Fraction >2U error: 4.29%

Proxy % hypo-risk: 2.86%

===== Safety Flags Summary =====

too_high_per_kg 74

low_time_gap 0

excessive_absolute 0

iob_exceeds 0

safety_flag 4

===== Calibration Table =====

| | bin | mean_pred_prob | observed_rate | count |
|---|-----------------|----------------|---------------|-------|
| 0 | (0.0825, 0.295] | 0.230021 | 1.000000 | 70 |
| 1 | (0.295, 0.561] | 0.375033 | 1.000000 | 70 |

✓ 8. Probabilistic / Quantile Regression

Train quantile models, produce median + lower/upper predictions, and show calibration & coverage checks.

Train quantile regressors

```
# Train quantile regressors (low / med / high)
```

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error
import numpy as np
import pandas as pd
```

```
# Use X (one-hot encoded) and y from earlier. If X is numpy array, convert to DataFrame:
if not isinstance(X, pd.DataFrame):
    X = pd.DataFrame(X, columns=available_features) # available_features from your prep
```

```
# Train/val split (reset indices)
Xtr, Xv, ytr, yv = train_test_split(X, y, test_size=0.2, random_state=42)
Xtr = Xtr.reset_index(drop=True); Xv = Xv.reset_index(drop=True)
ytr = np.array(ytr); yv = np.array(yv)
```

```
# Define quantile models
q_low = GradientBoostingRegressor(loss='quantile', alpha=0.10, n_estimators=200, max_depth=3,
q_med = GradientBoostingRegressor(loss='squared_error', n_estimators=200, max_depth=3, ra
q_high = GradientBoostingRegressor(loss='quantile', alpha=0.90, n_estimators=200, max_dep
```

```
# Fit models
q_low.fit(Xtr, ytr)
q_med.fit(Xtr, ytr)
q_high.fit(Xtr, ytr)

# Predict on validation
pred_low = q_low.predict(Xv)
pred_med = q_med.predict(Xv)
pred_high = q_high.predict(Xv)

# Quick metrics for median
mae_med = mean_absolute_error(yv, pred_med)
rmse_med = np.sqrt(((yv - pred_med)**2).mean())
```

```
print(f"Quantile models trained – validation MAE (median): {mae_med:.4f}, RMSE: {rmse_med
```

```
Quantile models trained – validation MAE (median): 0.6115, RMSE: 1.0835
```

Interval coverage, calibration, and visualization

#Coverage

```
import matplotlib.pyplot as plt
import seaborn as sns

# Coverage: fraction of true values inside [low, high]
inside = ((yv >= pred_low) & (yv <= pred_high)).astype(int)
coverage = inside.mean()
print(f"Interval coverage (p10-p90): {coverage:.3f}")

# Interval width statistics
width = pred_high - pred_low
print(f"Median interval width: {np.median(width):.3f}, mean width: {np.mean(width):.3f}")

Interval coverage (p10-p90): 0.629
Median interval width: 1.085, mean width: 1.572
```

#Calibration

```
def estimate_hypo_prob_quantile(y_pred_low, y_pred_med, y_pred_high, glucose_levels, isf_
"""
Monte Carlo sampling using Normal approx based on quantiles to estimate P(glucose_aft
glucose_after = glucose - sampled_bolus * ISF
"""

# Estimate std from 10-90 span. Use z factor ~2.5631 (see earlier cell)
denom = 2.5631
mu = y_pred_med
sigma = np.maximum((y_pred_high - y_pred_low) / denom, 1e-3)
probs = []
for i in range(len(mu)):
    samples = np.random.normal(mu[i], sigma[i], size=n_samples)
    samples = np.clip(samples, 0, None)
```

```

    ISF = 50.0 if isf_series is None else isf_series[i] # fallback
    pred_gluc_after = glucose_levels[i] - samples * ISF
    probs.append(np.mean(pred_gluc_after < 70))
return np.array(probs)

# If you have glucose_level and ISF aligned with Xv in a test_df, use them:
# Create test_df corresponding to Xv using model_df indices
test_df = model_df.loc[Xv.index].reset_index(drop=True).copy() # model_df was used earlie
gluc_levels = test_df['glucose_level'].values
isf_vals = test_df.get('median_ISF', test_df.get('correction_factor', pd.Series(50.0, ind

pred_hypo_prob = estimate_hypo_prob_quantile(pred_low, pred_med, pred_high, gluc_levels,

# calibration table: bucket by predicted prob and compare to naive observed (proxy)

# Make a fresh DataFrame aligned to validation set
test_df = pd.DataFrame(Xv, columns=X.columns).reset_index(drop=True)

# Add the true target (yv) and predictions
test_df['true'] = yv.reset_index(drop=True) if hasattr(yv, "reset_index") else yv
test_df['pred_low'] = pred_low
test_df['pred_med'] = pred_med
test_df['pred_high'] = pred_high
test_df['pred_hypo_prob'] = pred_hypo_prob

import pandas as pd

# Convert yv to Series so it aligns nicely
test_df['recommended_bolus_units'] = pd.Series(yv).reset_index(drop=True)

# Align ISF values from validation rows only
isf_vals = model_df.loc[Xv.index, 'median_ISF'].reset_index(drop=True)

# Hypoglycemia proxy flag
test_df['observed_hypo_proxy'] = (
    test_df['glucose_level'] - test_df['recommended_bolus_units'] * isf_vals
) < 70

# Calibration bins and table

test_df['prob_bin'] = pd.qcut(test_df['pred_hypo_prob'], q=5, duplicates='drop')
calib = test_df.groupby('prob_bin', observed=True).agg(
    mean_pred_prob=('pred_hypo_prob', 'mean'),
    observed_rate=('observed_hypo_proxy', 'mean'),
    count=('pred_hypo_prob', 'count')
).reset_index()

print("\nCalibration table (binned):")
display(calib)

```



Calibration table (binned):

| | prob_bin | mean_pred_prob | observed_rate | count | |
|---|-----------------|----------------|---------------|-------|--|
| 0 | (-0.001, 0.287] | 0.022024 | 0.071429 | 56 | |
| 1 | (0.287, 0.993] | 0.821075 | 0.516129 | 31 | |
| 2 | (0.993, 1.0] | 0.999811 | 0.849057 | 53 | |

Next steps:

[Generate code with calib](#)


[View recommended plots](#)

[New interactive sheet](#)

Plot

#calibration curve

```
plt.figure(figsize=(6,5))
plt.plot(calib['mean_pred_prob'], calib['observed_rate'], marker='o')
plt.plot([0,1],[0,1], '--', color='gray')
plt.xlabel("Mean predicted hypo probability")
plt.ylabel("Observed hypo rate (proxy)")
plt.title("Calibration curve (proxy)")
plt.grid(True)
plt.show()
```

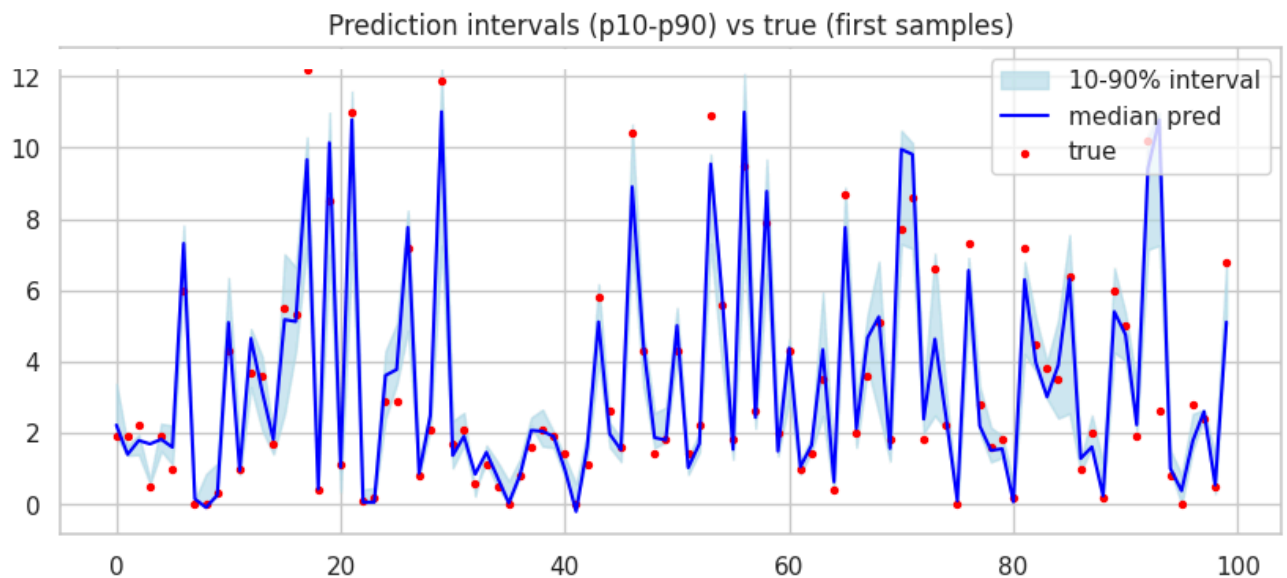


#Example prediction intervals vs true values (first 100 samples)

```

n_plot = min(100, len(pred_med))
plt.figure(figsize=(10,4))
x = np.arange(n_plot)
plt.fill_between(x, pred_low[:n_plot], pred_high[:n_plot], color='lightblue', alpha=0.6,
plt.plot(x, pred_med[:n_plot], label='median pred', color='blue')
plt.scatter(x, yv[:n_plot], color='red', s=8, label='true')
plt.legend()
plt.title("Prediction intervals (p10-p90) vs true (first samples)")
plt.show()

```



Simulator integration

This simulator is a simple deterministic compartment model that:

Models carb absorption (mono-exponential Ra),

Models insulin action as a decaying effect from bolus + basal,

Integrates glucose differential equation with Euler method (dt in minutes). It's not UVa/Padova but useful for evaluation and development — swap with a real simulator later.

#Minimal glucose-insulin simulator

```

import numpy as np
import pandas as pd

```

```

def simulate_minimal_model(row, bolus_units, horizon_minutes=240, dt_min=5,
                            Vg=10.0,      # glucose distribution volume (L) - heuristic
                            SI=0.02,      # insulin sensitivity (mg/dL per U per unit action
                            carb_abs_tau_min=None, # carbs absorption time constant
                            DIA_hours=5.0):
    """
    Simulate glucose trajectory for a single timepoint row given a bolus insulin (units).
    Returns a DataFrame with columns ['t_min', 'glucose'] and summary metrics.
    - row: pd.Series (should contain glucose_level, carbs_g, weight_kg, median_ISF option
    - bolus_units: units of insulin given at t=0
    """

```


Note: this is a simplified model for evaluation only.

"""

Initial / patient vars

G0 = float(row.get('glucose_level', 120.0)) # mg/dL

carbs = float(row.get('carbs_g', 0.0)) # grams

weight = float(row.get('weight_kg', 70.0))

carb absorption time constant (minutes)

if carb_abs_tau_min is None:

choose based on meal type or carb_duration if available

if 'carb_duration_min' in row and not pd.isna(row['carb_duration_min']):

carb_abs_tau_min = row['carb_duration_min'] / 2.0 + 20.0

else:

carb_abs_tau_min = 60.0

insulin action decay time constant (minutes) from DIA

tau_ins = DIA_hours * 60.0 / 2.0 # simple mapping: tau ~ DIA/2 (tunable)

time grid

times = np.arange(0, horizon_minutes + dt_min, dt_min) # minutes

n = len(times)

Initialize state variables

G = np.zeros(n); G[0] = G0

Ra = np.zeros(n) # rate of appearance (mg/dL per min)

Ieff = np.zeros(n) # insulin action (arbitrary units)

convert carbs (g) to mg/dL change: 1 g carbs ≈ 4 kcal -> approx glucose mg/dL chang

Use scaling: assume 1 g carbs raises blood glucose by approx 3 mg/dL / (Vg factor)

carb_to_gluc_factor = 3.0 # tweakable

Bolus generates initial Ieff impulse

Ieff[0] = bolus_units

simulate with Euler integration

for t_idx in range(1, n):

time step

dt = dt_min

carbohydrate absorption: monoexponential Ra (grams -> glucose units)

Ra_g/min = carbs / tau * exp(-t/tau)

t = times[t_idx]

Ra_g = (carbs / max(1.0, carb_abs_tau_min)) * np.exp(-t / carb_abs_tau_min) # g

Ra_gluc = Ra_g * carb_to_gluc_factor # convert to mg/dL per min (approx)

Ra[t_idx] = Ra_gluc

insulin action decay (exponential)

Ieff[t_idx] = Ieff[t_idx-1] * np.exp(-dt / tau_ins)

net glucose change: appearance - insulin-mediated uptake (proportional to Ieff*

insulin_term = SI * Ieff[t_idx] * max(0, G[t_idx-1] - 80.0) # uptake proportiona

dG = (Ra_gluc - insulin_term) * (dt / 1.0) / Vg # simplified scaling

small basal drift towards 100 mg/dL

basal_drift = 0.001 * (100.0 - G[t_idx-1]) * dt

G[t_idx] = max(30.0, G[t_idx-1] + dG + basal_drift)

sim_df = pd.DataFrame({'t_min': times, 'glucose': G, 'Ra': Ra, 'Ieff': Ieff})

summary metrics

```

min_gluc = float(G.min())
max_gluc = float(G.max())
time_in_range = float(((G >= 70.0) & (G <= 180.0)).mean() * 100.0)
hypo_episodes = int(((G < 70.0).astype(int)).sum())
return {'sim_df': sim_df, 'min_gluc': min_gluc, 'max_gluc': max_gluc,
        'time_in_range_pct': time_in_range, 'hypo_counts': hypo_episodes}

# Apply simulator to validation set predictions

# Inputs expected: test_df (aligned with Xv), pred_med (median predictions array), option

# Build test_df if not available:
test_df = model_df.loc[Xv.index].reset_index(drop=True).copy() # ensures alignment

# We'll evaluate the median predictions pred_med
results_sim = []
for i in range(len(pred_med)):
    row = test_df.iloc[i]
    pred_bolus = float(pred_med[i])
    sim_res = simulate_minimal_model(row, bolus_units=pred_bolus, horizon_minutes=240, dt
    # gather
    results_sim.append({
        'idx': i,
        'pred_bolus': pred_bolus,
        'true_bolus': float(row['recommended_bolus_units']),
        'min_gluc': sim_res['min_gluc'],
        'max_gluc': sim_res['max_gluc'],
        'time_in_range_pct': sim_res['time_in_range_pct'],
        'hypo_counts': sim_res['hypo_counts']
    })

sim_results_df = pd.DataFrame(results_sim)

# Summary statistics
print("Simulator-based summary (median predictions):")
print("Mean TIR%:", sim_results_df['time_in_range_pct'].mean())
print("Fraction with any hypo episodes:", (sim_results_df['hypo_counts'] > 0).mean())
print(sim_results_df.describe()[['min_gluc', 'time_in_range_pct', 'hypo_counts']])

➡ Simulator-based summary (median predictions):
Mean TIR%: 83.23615160349853
Fraction with any hypo episodes: 0.0

```

| | min_gluc | time_in_range_pct | hypo_counts |
|-------|------------|-------------------|-------------|
| count | 140.000000 | 140.000000 | 140.0 |
| mean | 120.739254 | 83.236152 | 0.0 |
| std | 33.288206 | 29.624636 | 0.0 |
| min | 70.300000 | 0.000000 | 0.0 |
| 25% | 92.389330 | 82.653061 | 0.0 |
| 50% | 112.428365 | 100.000000 | 0.0 |
| 75% | 144.589059 | 100.000000 | 0.0 |
| max | 204.966998 | 100.000000 | 0.0 |

```

# Example: show some rows where hypo occurred
print("\nExample rows with hypo episodes (sim):")

```

```
display(sim_results_df[sim_results_df['hypo_counts']>0].head())
```



Example rows with hypo episodes (sim):

| idx | pred_bolus | true_bolus | min_gluc | max_gluc | time_in_range_pct | hypo_counts |
|-----|------------|------------|----------|----------|-------------------|-------------|
|-----|------------|------------|----------|----------|-------------------|-------------|

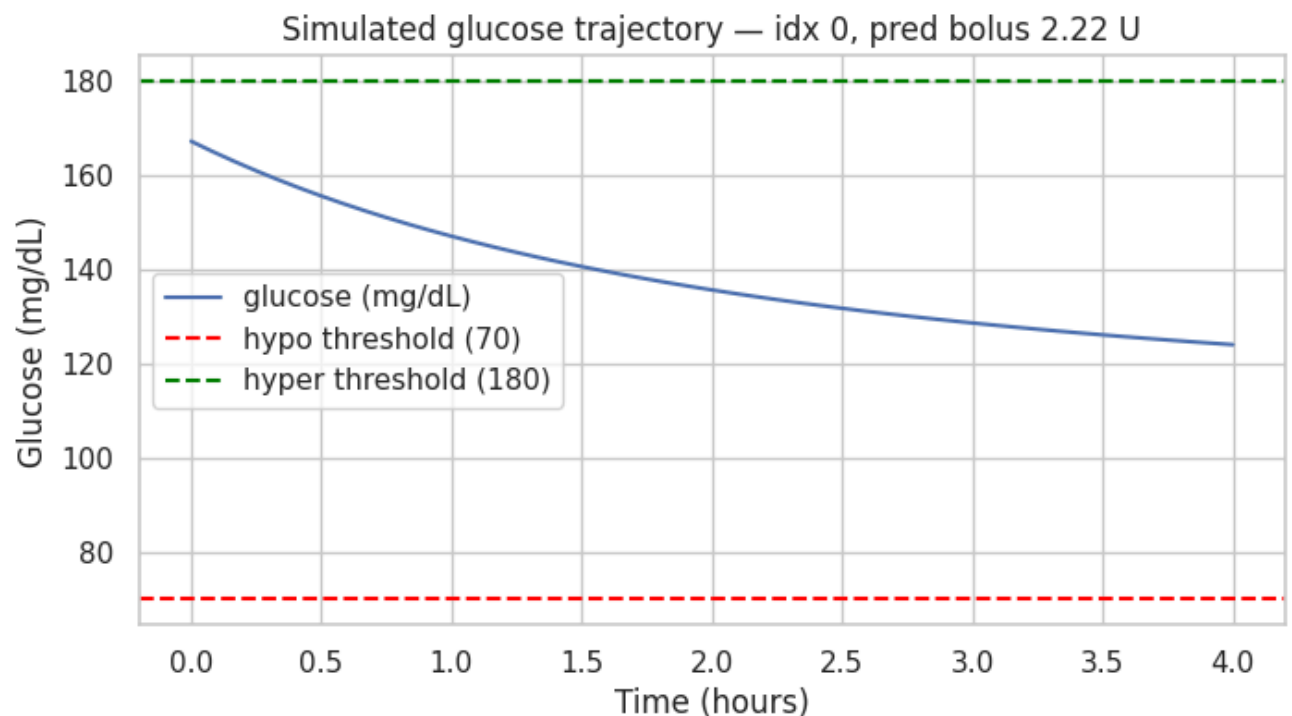


```
#Plot simulated glucose trajectory
```

```
import matplotlib.pyplot as plt
```

```
i_plot = 0 # change index as needed, must be < len(pred_med)
row = test_df.iloc[i_plot]
pred_bolus = float(pred_med[i_plot])
sim_res = simulate_minimal_model(row, bolus_units=pred_bolus, horizon_minutes=240, dt_min
sim_df = sim_res['sim_df']
```

```
plt.figure(figsize=(8,4))
plt.plot(sim_df['t_min']/60.0, sim_df['glucose'], label='glucose (mg/dL)')
plt.axhline(70, color='red', linestyle='--', label='hypo threshold (70)')
plt.axhline(180, color='green', linestyle='--', label='hyper threshold (180)')
plt.xlabel('Time (hours)')
plt.ylabel('Glucose (mg/dL)')
plt.title(f"Simulated glucose trajectory – idx {i_plot}, pred bolus {pred_bolus:.2f} U")
plt.legend()
plt.show()
```



✓ 10. Final Predict

```
from sklearn.model_selection import train_test_split  
from sklearn.metrics import mean_absolute_error, mean_squared_error
```